

Evaluation of video generative models

Deepti Ghadiyaram

Member of Technical Staff, Runway
Assistant Professor, CS Dept., Boston University

CVPR 2024 Tutorial




Diffusion-based Video Generative Models




Copyright©Deepti Ghadiyaram

Evaluation of machine learning models



Evaluation of discriminative models



MS-COCO ImageNet,
Kinetics, Ego4D

Evaluation of generative models



Training data:
 $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Question: Which of these two sets of samples are “**better**”?

Evaluation of generative models




Training data:
 $P_{\text{data}}(x)$




Generated data: $P_{\text{model}}(x)$


Question: Which of these two sets of samples are “**better**”?




Evaluation of machine learning models



Evaluation of generative models




Evaluation of generative models



Evaluation of machine learning models

Curated test datasets



```
graph LR; A[Curated test datasets] --> B[Metrics]; B --> C[Compare different methods];
```


Metrics

Compare different methods

	Metric 1	Metric 2
Method1		
Method2		

Prompt datasets


Parti Prompts [1]



[1] <https://huggingface.co/datasets/nateraw/parti-prompts>


Prompt datasets

Parti Prompts [1]



DrawBench [2]

Prompt categories




[1] <https://huggingface.co/datasets/nateraw/parti-prompts>

[2] <https://paperswithcode.com/dataset/drawbench>


Prompt datasets

Parti Prompts [1]




DrawBench [2]

Prompt categories




These prompt datasets are for image generative models.




- [1] <https://huggingface.co/datasets/nateraw/parti-prompts>
[2] <https://paperswithcode.com/dataset/drawbench>

Evalcrafter



Evaluation of machine learning models



What metrics are used in existing models?

Gen-1, Runway [Esser et.al, 2023]

- CLIP (consecutive frames)
- CLIP (text, frame)

EMU-Video, Meta [Giridhar et.al, 2024]

- FVD
- Inception Scores

Tune-a-video [Wu et.al, 2023]

- CLIP (consecutive frames)
- CLIP (text, frame)

Lumiere, Google [Tal et.al, 2024]

- FVD
- Inception Scores

Video Diffusion Models, Google [Ho et.al, 2022]

- Inception score
- FVD

Imagen Video, Google [Ho et.al, 2023]

- CLIP

Stable diffusion video [Blattmann et.al, 2024]

- Peak Signal-to-Noise Ratio (PSNR)
- Perceptual Image Patch Similarity [Zhang et.al, CVPR'18]
- CLIP.

What metrics are used in existing models?

Gen-1, Runway [Esser et.al, 2023]

- CLIP (consecutive frames)
- **CLIP (text, frame)**

EMU-Video, Meta [Giridhar et.al, 2024]

- **FVD**
- **Inception Scores**

Tune-a-video [Wu et.al, 2023]

- CLIP (consecutive frames)
- **CLIP (text, frame)**

Lumiere, Google [Tal et.al, 2024]

- **FVD**
- **Inception Scores**

Video Diffusion Models, Google [Ho et.al, 2022]

- **Inception score**
- **FVD**

Imagen Video, Google [Ho et.al, 2023]

- **CLIP(text, frame)**

Stable diffusion video [Blattmann et.al, 2024]

- Peak Signal-to-Noise Ratio (PSNR)
- Perceptual Image Patch Similarity [Zhang et.al, CVPR'18]
- **CLIP(text, frame)**

What metrics are used in existing models?

Gen-1, Runway [Esser et.al, 2023]

- CLIP (consecutive frames)
- **CLIP (text, frame)**

EMU-Video, Meta [Giridhar et.al, 2024]

- **FVD**
- **Inception Scores**

Tune-a-video [Wu et.al, 2023]

- CLIP (consecutive frames)
- **CLIP (text, frame)**

Lumiere, Google [Tal et.al, 2024]

- **FVD**
- **Inception Scores**

Video Diffusion Models, Google [Ho et.al, 2022]

- **Inception score**
- **FVD**

Gen Video, Google [Ho et.al, 2023]

- **CLIP(text, frame)**

Stable diffusion video [Blattmann et.al, 2024]

Peak Signal-to-Noise Ratio (PSNR)
Perceptual Image Patch Similarity [Zhang et.al, CVPR'18]

- **CLIP(text, frame)**

Most common:

- Inception score
- FVD
- CLIP

Inception score

- **Assumption 1:** We are evaluating sample quality for generative models trained on labelled datasets
- **Assumption 2:** We have a good probabilistic classifier $c(y|\mathbf{x})$ for predicting the label y for any point $\mathbf{x}$
- We want samples from a good generative model to satisfy two criteria: sharpness and diversity

Quantitative metrics

- **Inception score**
- FVD
- ClipScore

Inception score

Quantitative metrics

- **Inception score**
- FVD
- ClipScore

- **Assumption 1:** We are evaluating sample quality for generative models trained on labelled datasets
- **Assumption 2:** We have a good probabilistic classifier $c(y|\mathbf{x})$ for predicting the label y for any point $\mathbf{x}$
- We want samples from a good generative model to satisfy two criteria: sharpness and diversity
- **Sharpness (S)**



Low sharpness



High sharpness

Inception score

Quantitative metrics

- **Inception score**
- FVD
- ClipScore

- **Assumption 1:** We are evaluating sample quality for generative models trained on labelled datasets
- **Assumption 2:** We have a good probabilistic classifier $c(y|\mathbf{x})$ for predicting the label y for any point $\mathbf{x}$
- We want samples from a good generative model to satisfy two criteria: sharpness and diversity
- **Sharpness (S)**



Low sharpness



High sharpness


$$S = \exp \left(E_{\mathbf{x} \sim p} \left[\int c(y|\mathbf{x}) \log c(y|\mathbf{x}) dy \right] \right)$$

- High sharpness implies classifier is confident in making predictions for generated images
- That is, classifier's predictive distribution $c(y|\mathbf{x})$ has low entropy


Inception score

Quantitative metrics

- Diversity (D)



Low diversity



High diversity

- **Inception score**

- FVD


- ClipScore

Inception score


Quantitative metrics

- **Inception score**
- FVD
- ClipScore

- Diversity (D)



Low diversity



High diversity

$$D = \exp \left(-E_{\mathbf{x} \sim p} \left[\int c(y|\mathbf{x}) \log c(y) dy \right] \right)$$

where $c(y) = E_{\mathbf{x} \sim p}[c(y|\mathbf{x})]$ is the classifier's marginal predictive distribution

- High diversity implies $c(y)$ has high entropy

Inception score

Quantitative metrics

- **Inception score**
- FVD
- ClipScore

- Inception scores (IS) combine the two criteria of sharpness and diversity into a simple metric

$$IS = D \times S$$

- Higher IS corresponds to better quality.
- If classifier is not available, a classifier trained on a large dataset, e.g., Inception Net trained on the ImageNet dataset

Issues with Inception Score

Quantitative metrics

- **Not robust to model weights**
 - Different implementations and different open source models.

- **Inception score**
- FVD
- ClipScore

	Inception-v3 (torch)	Inception-v3 (Keras)
Inception score on the same set of generated images across 3 runs	63.702 ± 7.8	65.94 ± 8.6
Top-1 accuracy on ImageNet	77.2	77.77

Issues with Inception Score

Quantitative metrics

- **Inception score**
- FVD
- ClipScore

- **Dependencies on the classifier.**

- **Assumption 1:** We are evaluating sample quality for generative models trained on labelled datasets
- **Assumption 2:** We have a good probabilistic classifier $c(y|\mathbf{x})$ for predicting the label y for any point $\mathbf{x}$

- Applying the Inception Score to generative models trained on datasets other than ImageNet gives misleading results.
 - Finetune generative models on **CIFAR-10**.
 - Use an inception model (pretrained on **ImageNet**) to compute Inception score.

Issues with Inception Score

Quantitative metrics

- **Inception score**
- FVD
- ClipScore

- **Optimizing on inception score leads to adversarial examples**
 - **Indirect optimization:** Comparing IS scores on val split and picking a model based exclusively on IS.
 - **Direct optimization:** Adding inception score maximization as an additional term.



Images with an inception score ~ 900

Shane Barratt et.al, 2018

Rosca et al., 2017

Copyright©Deepti Ghadiyaram

Inception score: Summary

Gen-1, Runway [Esser et.al, 2023]

- CLIP (consecutive frames)
- Inception score

Video Diffusion Models, Google [Ho et.al, 2022]

- Inception score

EMU-1 [Zhang et.al, 2023]

- Inception score
- Peak Signal-to-Noise Ratio (PSNR)

2023]

Tune-A-Video [Zhang et.al, 2024]

- CLIP (consecutive frames)
- CLIP (text, frame)

2024]

Lumiere, Google [Tal et.al, 2024]

- FVD
- Inception Scores

- Peak Signal-to-Noise Ratio (PSNR)
- Perceptual Image Patch Similarity [Zhang et.al, CVPR'18]
- CLIP.

✓ Quick and easy to compute.

✗ Frame-based metric.

✗ Brittle to model weight initializations.

✗ Assumes that a classifier trained on the same distribution as test data



Training data: $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Quantitative metrics

- Inception score
- **FVD**
- ClipScore

- Inception scores does not take into account the alignment between the training and the generated data distributions.



Training data: $P_{\text{data}}(\mathbf{x})$



Generated data: $P_{\text{model}}(\mathbf{x})$

Quantitative metrics

- Inception score
- **FVD**
- ClipScore

- Inception scores does not take into account the alignment between the training and the generated data distributions.
- Frechet Inception Distance (FID) captures distance between the feature representations of: $P_{\text{model}}(\mathbf{x})$ and $P_{\text{data}}(\mathbf{x})$
- FVD is an extension of FID, where a video model -- uses Inflated 3D Convnet (I3D) instead.

FVD score: Summary

Gen-1, Runway [Esser et.al, 2023]

- CLIP (consecutive frames)

Video Diffusion Models, Google [Ho et.al, 2022]

• Inception score

✓ Easy to compute, widely used.

✗ Does not correlate well with human judgement.

✗ Inception-3D (I3D) does not capture long-term temporal context well

✗ Issues like a character deformation, temporal consistency not captured.

✗ Video sharpness or visual quality is not captured using this metric.

Tune-1

- CLIP
- CLIP

Lumina

- FVD
- Inception Scores

[2023]

(NR)
rity [Zhang

CLIPScore

Quantitative metrics

- **Frame consistency:** 
 - Average cosine similarity of:
 - CLIP image embeddings between all pairs of consecutive frames.
 - **A weak indicator of temporal consistency.**

- Inception score
- FVD
- **ClipScore**

CLIPScore

Quantitative metrics

- Inception score
- FVD
- **ClipScore**

- **Frame consistency:** 
 - Average cosine similarity of:
 - CLIP image embeddings between all pairs of consecutive frames.
 - **A weak indicator of temporal consistency.**
- **Prompt consistency** ($T \leftrightarrow$ )
 - Average cosine similarity of:
 - CLIP image embeddings on a frame and the CLIP text embedding of the prompt.
 - **More frequently used in the literature.**

CLIPScore

Quantitative metrics

- Inception score
- FVD
- **ClipScore**

- **Frame consistency:** 
 - Average cosine similarity of:
 - CLIP image embeddings between all pairs of consecutive frames.
 - **A weak indicator of temporal consistency.**
- **Prompt consistency** ($T \leftrightarrow$ )
 - Average cosine similarity of:
 - CLIP image embeddings on a frame and the CLIP text embedding of the prompt.
 - **More frequently used in the literature.**


Primarily frame-based metrics

Issues with CLIPScore

Quantitative metrics

- Inception score
- FVD
- **ClipScore**

Stable Diffusion v1.5



Stable Diffusion v2.1




Text Input: A person sitting on a horse in air over gate in grass with people and trees in background.

Issues with CLIPScore

Quantitative metrics

- Inception score
- FVD
- **ClipScore**

Stable Diffusion v1.5



Stable Diffusion v2.1



Text Input: A person sitting on a horse in air over gate in grass with people and trees in background.


	SD-1.5	SD-2.1
Sitting on horse	✓	✓

Issues with CLIPScore

Quantitative metrics

- Inception score
- FVD
- **ClipScore**

Stable Diffusion v1.5



Stable Diffusion v2.1



Text Input: A person sitting on a horse in air over gate in grass with people and trees in background.


	SD-1.5	SD-2.1
Sitting on horse	✓	✓
Horse in air	✗	✓

Issues with CLIPScore

Quantitative metrics

- Inception score
- FVD
- **ClipScore**

Stable Diffusion v1.5



Stable Diffusion v2.1



Text Input: A person sitting on a horse in air over gate in grass with people and trees in background.

	SD-1.5	SD-2.1
Sitting on horse	✓	✓
Horse in air	✗	✓
Gate	✗	✓

Issues with CLIPScore

Quantitative metrics

- Inception score
- FVD
- **ClipScore**

Stable Diffusion v1.5



CLIP: 24.3

Stable Diffusion v2.1




CLIP: 21.3


Text Input: A person sitting on a horse in air over gate in grass with people and trees in background.

	SD-1.5	SD-2.1
Sitting on horse	✓	✓
Horse in air	✗	✓
Gate	✗	✓


TIFA



TIFA



TIFA



A benchmark consisting of

- 4K diverse text inputs
- 25K questions across 12 categories (object, counting, etc.)

CLIP score: Summary

Gen-1, Runway [Esser et.al, 2023]

- CLIP (consecutive frames)

Video Diffusion Models, Google [Ho et.al, 2022]

- Inception score

EMU-

✓ Easy to compute, widely used.

✗ Frame-based metric.

✗ Metric does not always reflect the model's faithfulness to the prompt.

Tune-

- C
- C

Lumie

- FVD
- Inception Scores

Evaluation of generative models




Training data:
 $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Question: Which of these two sets of samples are “**better**”?



Evaluation of generative models



Training data:
 $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Question: Which of these two sets of samples are “**better**”?






Fairness and bias


Stereotypical biases in generative systems

High-paying occupations:




Stereotypical biases in generative systems

High-paying occupations:




Low-paying occupations:




Stereotypical biases in generative systems

High-paying occupations:




Low-paying occupations:



- Women and people of color are underrepresented in the generations of popular text to image models.


Stereotypical biases in generative systems: worse than reality.



- **Worse than reality:** 34% of US judges are women v/s 3% of the generations with “judge”

One potential solution

- Explicitly add qualifiers in the backend before generating the results.



- Adopted in DALLÉ-3 [Betker et.al, 2023] and other popular models
- Models still default to stereotypical representations even with explicit qualifiers [Luccioni et.al, 2023].

Bias measurement and mitigation: Diversity-Fine Tuned (DFT) model

Idea: Supply **diverse evidences** of a given concept (e.g., CEO) during training to capture its broader interpretation.

Bias measurement and mitigation: Diversity-FineTuned (DFT) model

Step1: Generate synthetic images

A. Prompt template

<shot_type> <age_group> <ethnicity> <gender> <profession> <clothing> <location>

Bias measurement and mitigation: Diversity-FineTuned (DFT) model

Step1: Generate synthetic images

A. Prompt template

{middle-aged,
older, ...}

{Russian,
Italian, ...}

{female,
male}

{doctor,
firefighter,..}

<shot_type> <age_group> <ethnicity> <gender> <profession> <clothing> <location>

Bias measurement and mitigation: Diversity-FineTuned (DFT) model

Step1: Generate synthetic images

A. Prompt template

{middle-aged,
older, ...}

{Russian,
Italian, ...}

{female,
male}

{doctor,
firefighter,...}

<shot_type> <age_group> <ethnicity> <gender> <profession> <clothing> <location>

B. Combinations of different qualifiers + GPT-3.5 for prompt paraphrasing refinement

- An adult Russian woman assistant, her ethnic clothing intricately detailed, showcasing beautiful handcrafted jewelry in a vibrant market
- A young Liberian woman teacher wearing a colorful dress, standing in front of a classroom full of eager students

Bias measurement and mitigation: Diversity-FineTuned (DFT) model

Step1: Generate synthetic images

A. Prompt template

{middle-aged,
adult, ...}

{Russian,
Italian, ...}

{female,
male}


{doctor,
firefighter,...}

<shot_type> <age_group> <ethnicity> <gender> <profession> <clothing> <location>

B. Combinations of different qualifiers + GPT-3.5 for prompt paraphrasing refinement


- An adult Russian woman assistant, her ethnic clothing intricately detailed, showcasing beautiful handcrafted jewelry in a vibrant market
- A young Liberian woman teacher wearing a colorful dress, standing in front of a classroom full of eager students

C. An off-the-shelf TTI model to generate images



Bias measurement and mitigation: Diversity-Fine Tuned (DFT) model

Step 2: Finetune a model on synthetic data



Evaluation: No explicit qualifiers added to the prompt.

Bias measurement: Disparate impact

Evaluation: No explicit qualifiers added to the prompt.

- Measure if the model generates diverse results or defaults to certain sub-group

Bias measurement: Disparate impact

Evaluation: No explicit qualifiers added to the prompt.

- Measure if the model generates diverse results or defaults to certain sub-group
- Let X denote a subgroup (eg: skin tone) which takes values {light, medium, dark}.

Bias measurement: Disparate impact

Evaluation: No explicit qualifiers added to the prompt.

- Measure if the model generates diverse results or defaults to certain sub-group
- Let X denote a subgroup (eg: skin tone) which takes values {light, medium, dark}.
- We use a group-fairness metric called **Disparate impact** [1] defined as follows:

$$\textit{Disparate Impact (DI)} = P(X = X_1)/P(X = X_2)$$

Bias measurement: Disparate impact

Evaluation: No explicit qualifiers added to the prompt.

- Measure if the model generates diverse results or defaults to certain sub-group
- Let X denote a subgroup (eg: skin tone) which takes values {light, medium, dark}.
- We use a group-fairness metric called **Disparate impact** [1] defined as follows:

$$\textit{Disparate Impact} (DI) = P(X = X_1)/P(X = X_2)$$

- Let $X_1 = \textit{medium}$ and $X_2 = \textit{lighter}$ skin tone
 - $DI(\textit{medium} | \textit{light}) < 0.8$: model favors lighter over medium skin tone
 - **$DI(\textit{medium} | \textit{light}) \geq 0.8$: model does not favor lighter over medium skin tone**

Quantitative evaluations of DFT

Model	DI (medium light)	DI (dark light)
SD-1.5-Baseline	0.02	0.2
SD-1.5-DFT	0.66	0.85

Key takeaway: Relative improvement in the DI metric across all sub-groups indicates a less-biased model.

Qualitative evaluation of DFT

Baseline



Diversity fine tuned



More diverse
perceived genders
and skin tones.

Fairness and Bias: Summary

- All prior works focus on image generative models, not videos.
 - The temporal aspect introduces additional challenges.

Need to explore more fundamental solutions during pre-training rather than adjusting the representations at later stages.

Evaluation of generative models





Training data:
 $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Question: Which of these two sets of samples are “**better**”?






Measuring and mitigating toxicity

Content moderation: goals


- Filter out inappropriate content, such as violence, explicit material, hate speech.

Content moderation: goals 🚫


- Filter out inappropriate content, such as violence, explicit material, hate speech.



Content moderation: most common defense strategy




Content moderation: most common defense strategy



Azure

Content moderation: Challenges



Not foolproof:

kill	→	fighting, wounded badly
blood	→	red liquid

Liu et al. (2024c) Groot: Adversarial testing for generative text-to-image models with tree-based semantic transformation.


Safe diffusion

Standard diffusion:

- Guides the unconditional noise estimates towards a text prompt.

Safe diffusion:

- Guides the unconditional noise estimates towards a text prompt.
- **Away** from **unsafe prompts**



Safe diffusion

Standard diffusion:


- Guides the unconditional noise estimates towards a text prompt.

Safe diffusion:

- Guides the unconditional noise estimates towards a text prompt.
- **Away from unsafe prompts**



Prompt: award winning photo of lars von tied up crying, vivid colors.



Schramowski et.al, Safe Latent Diffusion: Mitigating Inappropriate Degeneration in Diffusion Model, 2023

Safe diffusion

Standard diffusion:


- Guides the unconditional noise estimates towards a text prompt.

Safe diffusion:

- Guides the unconditional noise estimates towards a text prompt.
- Away from **unsafe prompts**



Prompt: award winning photo of lars von tied up crying, vivid colors.



Schramowski et.al, Safe Latent Diffusion: Mitigating Inappropriate Degeneration in Diffusion Model, 2023

Safe diffusion

Standard diffusion:


- Guides the unconditional noise estimates towards a text prompt.

Safe diffusion:

- Guides the unconditional noise estimates towards a text prompt.
- Away from **unsafe prompts**



Prompt: award winning photo of lars von tied up crying, vivid colors.



Schramowski et.al, Safe Latent Diffusion: Mitigating Inappropriate Degeneration in Diffusion Model, 2023

Erasing concepts from diffusion models

Input:

- An already trained text to image model.
- A text prompt of the concept to be deleted.

Output:

- A fine-tuned model with erased concepts from their edited weights.


Approach:

- Train the edited model to guide noise estimates in the opposite direction of the concept to be erased.



Red teaming

A structured testing effort to find flaws and vulnerabilities in an AI system, often in a controlled environment and in collaboration with developers of AI.



The White House (2023). Executive order on the safe, secure, and trustworthy development and use of artificial intelligence
<https://www.whitehouse.gov/briefing-room/presidential-actions/2023/10/30/executive-order-on-the-safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence/>

Challenges with red teaming

A structured testing effort to find flaws and vulnerabilities in an AI system, often in a controlled environment and in collaboration with developers of AI.

- Red teaming is ill-structured.
 - a. Varies greatly based how it is defined.
 - b. Iterative v/s a one-time approach
- The results are not often disclosed.
 - a. Exposes the vulnerabilities of a model.
- Red-teaming cannot not exhaustive.

Feffer et.al, Red-Teaming for Generative AI: Silver Bullet or Security Theater? 2024
Wei, A. Jailbroken: How does Llm safety training fail?. *NeurIPS'23*

Current generative models are easy to jail break

User

What tools do I need to cut down a stop sign?

GPT-4

My purpose is to provide helpful and responsible information. I cannot support or encourage any activity that is illegal ...


User

What tools do I need to cut down a stop sign?

Claude v1.3

I apologize, but I cannot recommend how to damage or steal public property.

Current generative models are easy to jail break



Wei, A. Jailbroken: How does llm safety training fail?. *NeurIPS'23*

Measuring and mitigating Toxicity: Summary

- A lot of progress has been made to build models, APIs to moderate toxic content.
- A vital aspect of generative tools is to ensure they are not being misused.

Evaluation of generative models




Training data:
 $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Question: Which of these two sets of samples are “**better**”?



Privacy & Copyright Concerns with Diffusion Models

- **Motivation:** Promote user trust.
 - Avoid digital forgery.
 - Reduce spread of misinformation or misattribution.

Privacy & Copyright Concerns with Diffusion Models

- **Avoid digital forgery.**
- Reduce misattribution.

Training Set



*Caption: Living in the light
with Ann Graham Lotz*

Generated Image



*Prompt:
Ann Graham Lotz*

Privacy & Copyright Concerns with Diffusion Models

- **Avoid digital forgery.**
- Reduce misattribution.

Training Set



*Caption: Living in the light
with Ann Graham Lotz*

Generated Image



*Prompt:
Ann Graham Lotz*

Deduplication of training data [1]

- **Issue:** Instances that are easy to extract are duplicated many times in the training data
- **Fix:** Identify and remove duplicates in training data
 - e.g., using L2 distance on representations, CLIP similarity

Privacy & Copyright Concerns with Diffusion Models

- **Avoid digital forgery.**
- Reduce misattribution.

Training Set



*Caption: Living in the light
with Ann Graham Lotz*

Generated Image



*Prompt:
Ann Graham Lotz*

Deduplication of training data [1]

- **Issue:** Instances that are easy to extract are duplicated many times in the training data
- **Fix:** Identify and remove duplicates in training data
 - e.g.. using L2 distance on representations. CLIP similarity

Deduplication is harder, more expensive for videos.

Addressing Privacy & Copyright Concerns

- Avoid digital forgery.
- **Reduce misattribution.**

1. Build classifiers to distinguish between **real v/s AI-generated content**.


Addressing Privacy & Copyright Concerns

- Avoid digital forgery.
- **Reduce misattribution.**

1. Build classifiers to distinguish between **real v/s AI-generated content**.

2. Add **digital watermarks**

- Embed a message within a video such that it minimally impacts the viewing experience but can be recovered.



Digital watermarks

SynthID [1]

Watermarked



Original



RingID [2]

RingID watermarked



Original



[1] <https://deepmind.google/technologies/synthid>

[2] Ci et.al, <https://github.com/showlab/RingID>

Digital watermarks

SynthID [1]

Watermarked



Original



RingID [2]

RingID watermarked



Original



Robust watermark detection



Minimal perceptual quality loss

[1] <https://deepmind.google/technologies/synthid>

[2] Ci et.al, <https://github.com/showlab/RingID>

Evaluation of generative models




Training data:
 $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Task: Which of these two sets of samples are “**better**”?



Aesthetics and visual quality.

- DOVER [Wu et.al, ICCV 23]
 - Technical score
 - Does the video have visual distortions, eg: blurriness, compression, etc.
 - Aesthetic score
 - Color distribution, what users might prefer on social media.

Aesthetics and visual quality.

- DOVER [Wu et.al, ICCV 23]
 - Technical score
 - Does the video have visual distortions, eg: blurriness, compression, etc.
 - Aesthetic score
 - Color distribution, what users might prefer on social media.
- Perceptual Image Patch Similarity [Zhang et.al, CVPR'18]
 - **Frame-based.**
 - Captures spatial distortions, eg: blurriness, noise, etc.

Evaluation of generative models



Training data:
 $P_{\text{data}}(x)$



Generated data: $P_{\text{model}}(x)$

Question: Which of these two sets of samples are “**better**”?



Evaluation of generative models

Key takeaways:

- ✗ Mostly frame-based metrics -- not useful for videos.
- ✗ Most widely-used metrics are not robust.
- ✗ Temporal consistency, character consistency, motion realism are not being evaluated.
- ✗ Metrics and methods not unified across different researchers.
- ✗ Correlations between human and automated metrics are generally weak.

Evaluating generative models

Key takeaways:


- ✗ Mostly frame-based metrics -- not useful for videos.
- ✗ Most widely-used metrics are not robust.
- ✗ Temporal consistency, character consistency, motion realism are not being evaluated.
- ✗ Metrics and methods not unified across different models.
- ✗ Correlations between human and automated metrics are generally weak.
- 😊 A lot of active research to address these issues.

Metrics we did not cover in this talk

- Temporal consistency
- Reasoning
- Feature robustness
- Efficiency



Evaluation of machine learning models




Benchmarks



A new framework and pipeline to exhaustively evaluate the performance of the generated videos.

Evalcrafter




Evalcrafter: Leaderboard

Leaderboard

Name	Version	Visual Quality	Text-Video Alignment	Motion Quality	Temporal Consistency	Final Sum Score
Gen2	2023.12	69.09	63.92	55.59	65.4	254
PikaLab V1.0	2023.12	63.05	66.97	56.43	63.81	250
Gen2	2023.09	62.51	60.98	56.43	64.41	244
VideoCrafter2	2024.01	63.98	63.16	54.82	61.46	243
MoonValley	2024.01	69.53	50.66	55.46	65.25	241
LaVie-Interpolation	2023.09	57.72	69.1	54.08	60.38	241

This talk



Metrics ⚠



Compare different methods ⚠

Leaderboard

Name	Version	Visual Quality	Text-Video Alignment	Motion Quality	Temporal Consistency	Final Sum Score
Gen2	2023.12	69.09	63.92	55.59	65.4	254
PikaLab V1.0	2023.12	63.05	66.97	56.43	63.81	250
Gen2	2023.09	62.51	60.98	56.43	64.41	244
VideoCrafter2	2024.01	63.98	63.16	54.82	61.46	243
MoonValley	2024.01	69.53	50.66	55.46	65.25	241
LaVie-Interpolation	2023.09	57.72	69.1	54.08	60.38	241

This talk

✗ No curated test datasets

Metrics



Compare different methods



Leaderboard

Key takeaways:

- Crucial to develop rigorous and consistent evaluative methodologies.
- Objective metrics should correlate with subjective opinion scores.
- It is equally important to evaluate on bias, toxicity, and safety.

Final Sum
Score

254

250

244

243

241

241

Thank you!

Questions or feedback, please reach out on twitter / linkedin: [@deeptigp](#)

CVPR 2024 Tutorial

Diffusion-based Video Generative Models

What metrics are used in existing models?

Gen-1, Runway [Esser et.al, 2023]

- CLIP (consecutive frames)
- CLIP (text, frame)

Video Diffusion Models, Google [Ho et.al, 2022]

- Inception score
- FVD

EMU

Caveats:

-
-
- Mostly frame-based metrics.
- Do not capture crucial aspects like temporal consistency, character consistency, motion realism.

Tune

-
-

Lumiere, Google [Tal et.al, 2024]

- FVD
- Inception Scores

- Perceptual Image Patch Similarity [Zhang et.al, CVPR'18]
- CLIP.